

Assignment 6

Nikhil Vemula, Atharva Pore, Shreeja Deshpande, Ami Sharma

November 17, 2022

Question 10

Observations

1. With no index postgres is performing sequential scan using multiple workers to filter
2. Creating an index on **skill** column reduced the execution time by almost 50%
3. Didn't find any major difference in execution times between btree index and hash index.
4. But hash index took longer time to create on on column with more number of records

size n of relation PersonSkill	No index	With index
10^4	1.007	0.574
10^5	9.066	5.825
10^6	92.875	45.834
10^7	973.207	639.765

Query plans

For 10^4 records

No Index Query Plan

```
QUERY PLAN
-----
Seq Scan on personskill (cost=0.00..163.69 rows=1802 width=4) (actual time=0.012..1.045 rows=1802 loops=1)
  Filter: (skill = 'AI'::text)
  Rows Removed by Filter: 7293
Planning Time: 0.109 ms
Execution Time: 1.191 ms
(5 rows)
```

Btree Index Query Plan

```

QUERY PLAN
-----
Bitmap Heap Scan on personskill  (cost=22.25..94.78 rows=1802 width=4) (actual time=0.067..0.465 rows=1802 loops=1)
  Recheck Cond: (skill = 'AI'::text)
  Heap Blocks: exact=50
  -> Bitmap Index Scan on personskill_skill_idx  (cost=0.00..21.80 rows=1802 width=0) (actual time=0.054..0.054 rows=1802 loops=1)
        Index Cond: (skill = 'AI'::text)
Planning Time: 0.053 ms
Execution Time: 0.620 ms
(7 rows)

```

For 10^5 records

No Index Query Plan

```

QUERY PLAN
-----
Seq Scan on personskill  (cost=0.00..1621.14 rows=18123 width=4) (actual time=0.007..8.684 rows=18172 loops=1)
  Filter: (skill = 'AI'::text)
  Rows Removed by Filter: 72399
Planning Time: 0.061 ms
Execution Time: 9.847 ms
(5 rows)

```

Btree Index Query Plan

```

QUERY PLAN
-----
Bitmap Heap Scan on personskill  (cost=208.75..924.28 rows=18123 width=4) (actual time=0.611..4.082 rows=18172 loops=1)
  Recheck Cond: (skill = 'AI'::text)
  Heap Blocks: exact=489
  -> Bitmap Index Scan on skill_index  (cost=0.00..204.21 rows=18123 width=0) (actual time=0.546..0.546 rows=18172 loops=1)
        Index Cond: (skill = 'AI'::text)
Planning Time: 0.231 ms
Execution Time: 5.322 ms
(7 rows)

```

For 10^6 records

No index Query Plan

```

QUERY PLAN
-----
Seq Scan on personskill  (cost=0.00..16227.24 rows=183887 width=4) (actual time=0.025..81.109 rows=181384 loops=1)
  Filter: (skill = 'AI'::text)
  Rows Removed by Filter: 726355
Planning Time: 0.040 ms
Execution Time: 91.967 ms
(5 rows)

```

Btree index Query Plan

```

QUERY PLAN
-----
Bitmap Heap Scan on personskill  (cost=2031.29..9186.98 rows=181015 width=4) (actual time=6.280..38.996 rows=181384 loops=1)
  Recheck Cond: (skill = 'AI'::text)
  Heap Blocks: exact=4893
  -> Bitmap Index Scan on skill_index  (cost=0.00..1986.04 rows=181015 width=0) (actual time=5.554..5.555 rows=181384 loops=1)
        Index Cond: (skill = 'AI'::text)
Planning Time: 0.238 ms
Execution Time: 50.507 ms
(7 rows)

```

For 10^7 records

No index Query Plan

```
queryplan
-----
Seq Scan on personskill (cost=0.00..162223.31 rows=1834672 width=4) (actual time=0.019..795.445 rows=1811314 loops=1)
  Filter: (skill = 'AI'::text)
  Rows Removed by Filter: 7253413
  Planning Time: 0.113 ms
  Execution Time: 895.138 ms
(5 rows)
```

Btree index Query Plan

```
queryplan
-----
Bitmap Heap Scan on personskill (cost=20527.37..92377.13 rows=1834701 width=4) (actual time=77.331..499.683 rows=1811314 loops=1)
  Recheck Cond: (skill = 'AI'::text)
  Heap Blocks: exact=48916
  -> Bitmap Index Scan on skill_index (cost=0.00..20068.69 rows=1834701 width=0) (actual time=68.578..68.578 rows=1811314 loops=1)
    Index Cond: (skill = 'AI'::text)
  Planning Time: 0.286 ms
  Execution Time: 603.233 ms
(7 rows)
```

Question 11

Observations

1. As expected adding an btree index helped for range queries. But!
2. If the stats for a given range are high then postgres is performing a sequential scan though index is available.
3. For smaller range, the execution time is nearly constant even when the number of records increased
4. Interestingly larger range performs slightly better than medium range as it covers the whole table?

Small Range

size n of relation worksFor	No index	With index
10^4	0.961	0.020
10^5	8.45	0.021
10^6	37.561	0.022
10^7	279.965	0.021

Medium Range

size n of relation worksFor	No index	With index
10^4	1.726	1.552
10^5	24.065	23.626
10^6	216.149	203.814
10^7	1970.034	1953.22

Large Range

size n of relation <code>worksFor</code>	No index	With index
10^4	2.565	2.409
10^5	20.516	21.637
10^6	196.60	205.833
10^7	1807.410	1783.657

Query Plans

Small Range

For 10^4 records

No Index Query Plan

```
QUERY PLAN
-----
Seq Scan on worksfor  (cost=0.00..205.00 rows=1 width=4) (actual time=0.009..0.964 rows=1 loops=1)
  Filter: ((10000 <= salary) AND (salary <= 10000))
  Rows Removed by Filter: 9999
  Planning Time: 0.148 ms
  Execution Time: 0.976 ms
(5 rows)
```

With Index Query Plan

```
QUERY PLAN
-----
Index Scan using worksfor_salary_idx on worksfor  (cost=0.29..8.32 rows=2 width=4) (actual time=0.007..0.008 rows=1 loops=1)
  Index Cond: ((salary >= 10000) AND (salary <= 10000))
  Planning Time: 0.128 ms
  Execution Time: 0.019 ms
(4 rows)
```

For 10^5 records

No Index Query Plan

```
QUERY PLAN
-----
Seq Scan on worksfor  (cost=0.00..2041.00 rows=1 width=4) (actual time=0.008..8.500 rows=3 loops=1)
  Filter: ((10000 <= salary) AND (salary <= 10000))
  Rows Removed by Filter: 99997
  Planning Time: 0.071 ms
  Execution Time: 8.513 ms
(5 rows)
```

With Index Query Plan

```
QUERY PLAN
-----
Index Scan using worksfor_salary_idx on worksfor  (cost=0.29..8.33 rows=2 width=4) (actual time=0.007..0.008 rows=3 loops=1)
  Index Cond: ((salary >= 10000) AND (salary <= 10000))
  Planning Time: 0.453 ms
  Execution Time: 0.021 ms
(4 rows)
```

For 10⁶ records

No Index Query Plan

```
QUERY PLAN
-----
Gather  (cost=1000.00..13620.10 rows=1 width=8) (actual time=0.135..37.094 rows=4 loops=1)
  Workers Planned: 2
  Workers Launched: 2
  -> Parallel Seq Scan on worksfor  (cost=0.00..12620.00 rows=1 width=8) (actual time=19.437..31.110 rows=1 loops=3)
        Filter: ((100 <= salary) AND (salary <= 100))
        Rows Removed by Filter: 333332
  Planning Time: 0.075 ms
  Execution Time: 37.107 ms
(8 rows)
```

With Index Query Plan

```
QUERY PLAN
-----
Index Scan using worksfor_salary_idx on worksfor  (cost=0.42..8.46 rows=2 width=8) (actual time=0.014..0.016 rows=4 loops=1)
  Index Cond: ((salary >= 100) AND (salary <= 100))
  Planning Time: 0.291 ms
  Execution Time: 0.027 ms
(4 rows)
```

For 10⁷ records

No Index Query Plan

```
QUERY PLAN
-----
Gather  (cost=1000.00..127195.82 rows=1 width=8) (actual time=0.173..282.248 rows=1 loops=1)
  Workers Planned: 2
  Workers Launched: 2
  -> Parallel Seq Scan on worksfor  (cost=0.00..126195.72 rows=1 width=8) (actual time=182.311..275.322 rows=0 loops=3)
        Filter: ((100 <= salary) AND (salary <= 100))
        Rows Removed by Filter: 3333333
  Planning Time: 0.130 ms
  Execution Time: 282.266 ms
(8 rows)
```

With Index Query Plan

```
QUERY PLAN
-----
Index Scan using worksfor_salary_idx on worksfor  (cost=0.43..8.50 rows=3 width=8) (actual time=0.014..0.015 rows=1 loops=1)
  Index Cond: ((salary >= 100) AND (salary <= 100))
  Planning Time: 0.501 ms
  Execution Time: 0.027 ms
(4 rows)
```

Medium Range

For 10⁴ records

No Index Query Plan

```
QUERY PLAN
-----
Seq Scan on worksfor  (cost=0.00..205.00 rows=4961 width=4) (actual time=0.014..1.688 rows=4955 loops=1)
  Filter: ((10000 <= salary) AND (salary <= 50186992))
  Rows Removed by Filter: 5045
  Planning Time: 0.073 ms
  Execution Time: 2.182 ms
(5 rows)
```

With Index Query Plan

```
QUERY PLAN
-----
Index Scan using worksfor_salary_idx on worksfor (cost=0.29..186.53 rows=4962 width=4) (actual time=0.018..1.170 rows=4955 loops=1)
  Index Cond: ((salary >= 10000) AND (salary <= 50186992))
  Planning Time: 0.090 ms
  Execution Time: 1.531 ms
(4 rows)
```

For 10^5 records

No Index Query Plan

```
QUERY PLAN
-----
Seq Scan on worksfor (cost=0.00..2291.00 rows=33330 width=4) (actual time=0.008..19.839 rows=50095 loops=1)
  Filter: ((10000 <= salary) AND ((salary)::numeric <= 499212216.5))
  Rows Removed by Filter: 49905
  Planning Time: 0.038 ms
  Execution Time: 23.056 ms
(5 rows)
```

With Index Query Plan

```
QUERY PLAN
-----
Seq Scan on worksfor (cost=0.00..2291.00 rows=33333 width=4) (actual time=0.010..20.457 rows=50095 loops=1)
  Filter: ((10000 <= salary) AND ((salary)::numeric <= 499212216.5))
  Rows Removed by Filter: 49905
  Planning Time: 0.083 ms
  Execution Time: 23.758 ms
(5 rows)
```

For 10^6 records

No Index Query Plan

```
QUERY PLAN
-----
Seq Scan on worksfor (cost=0.00..23870.00 rows=333300 width=8) (actual time=0.010..179.810 rows=500077 loops=1)
  Filter: ((100 <= salary) AND ((salary)::numeric <= 49956718.44850))
  Rows Removed by Filter: 499923
  Planning Time: 0.048 ms
  Execution Time: 209.021 ms
(5 rows)
```

With Index Query Plan

```
QUERY PLAN
-----
Seq Scan on worksfor (cost=0.00..23870.00 rows=333333 width=8) (actual time=0.013..211.058 rows=500077 loops=1)
  Filter: ((100 <= salary) AND ((salary)::numeric <= 49956718.44850))
  Rows Removed by Filter: 499923
  Planning Time: 0.116 ms
  Execution Time: 240.352 ms
(5 rows)
```

For 10^7 records

No Index Query Plan

```
QUERY PLAN
-----
Seq Scan on worksfor (cost=0.00..238697.01 rows=3333038 width=8) (actual time=0.129..1717.734 rows=5000037 loops=1)
  Filter: ((100 <= salary) AND ((salary)::numeric <= 500082613.5282))
  Rows Removed by Filter: 4999963
  Planning Time: 0.072 ms
  Execution Time: 1985.666 ms
(5 rows)
```

With Index Query Plan

```
QUERY PLAN
-----
Seq Scan on worksfor (cost=0.00..238695.00 rows=3333333 width=8) (actual time=0.010..1737.548 rows=5000037 loops=1)
  Filter: ((100 <= salary) AND ((salary)::numeric <= 500082613.5282))
  Rows Removed by Filter: 4999963
  Planning Time: 0.081 ms
  Execution Time: 2017.289 ms
(5 rows)
```

Large Range

For 10^4 records

No Index Query Plan

```
QUERY PLAN
-----
Seq Scan on worksfor (cost=0.00..205.00 rows=9998 width=4) (actual time=0.010..1.891 rows=10000 loops=1)
  Filter: ((10000 <= salary) AND (salary <= 99980000))
  Planning Time: 0.052 ms
  Execution Time: 2.656 ms
(4 rows)
```

With Index Query Plan

```
QUERY PLAN
-----
Seq Scan on worksfor (cost=0.00..205.00 rows=10000 width=4) (actual time=0.013..1.734 rows=10000 loops=1)
  Filter: ((10000 <= salary) AND (salary <= 99980000))
  Planning Time: 0.211 ms
  Execution Time: 2.500 ms
(4 rows)
```

For 10^5 records

No Index Query Plan

```
QUERY PLAN
-----
Seq Scan on worksfor (cost=0.00..2041.00 rows=99980 width=4) (actual time=0.010..14.724 rows=100000 loops=1)
  Filter: ((10000 <= salary) AND (salary <= 1000000000))
  Planning Time: 0.143 ms
  Execution Time: 21.192 ms
(4 rows)
```

With Index Query Plan

```

QUERY PLAN
-----
Seq Scan on worksfor  (cost=0.00..2041.00 rows=100000 width=4) (actual time=0.011..14.475 rows=100000 loops=1)
  Filter: ((10000 <= salary) AND (salary <= 10000000000))
  Planning Time: 0.227 ms
  Execution Time: 20.928 ms
(4 rows)

```

For 10^6 records

No Index Query Plan

```

QUERY PLAN
-----
Seq Scan on worksfor  (cost=0.00..21370.00 rows=999800 width=8) (actual time=0.024..136.545 rows=1000000 loops=1)
  Filter: ((100 <= salary) AND (salary <= 1000000000))
  Planning Time: 0.129 ms
  Execution Time: 194.728 ms
(4 rows)

```

With Index Query Plan

```

QUERY PLAN
-----
Seq Scan on worksfor  (cost=0.00..21370.00 rows=1000000 width=8) (actual time=0.009..135.051 rows=1000000 loops=1)
  Filter: ((100 <= salary) AND (salary <= 1000000000))
  Planning Time: 0.213 ms
  Execution Time: 192.895 ms
(4 rows)

```

For 10^7 records

No Index Query Plan

```

QUERY PLAN
-----
Seq Scan on worksfor  (cost=0.00..213696.73 rows=9998115 width=8) (actual time=0.030..1298.063 rows=10000000 loops=1)
  Filter: ((100 <= salary) AND (salary <= 10000000000))
  Planning Time: 0.064 ms
  Execution Time: 1837.039 ms
(4 rows)

```

With Index Query Plan

```

QUERY PLAN
-----
Seq Scan on worksfor  (cost=0.00..213695.00 rows=10000000 width=8) (actual time=0.010..1345.495 rows=10000000 loops=1)
  Filter: ((100 <= salary) AND (salary <= 10000000000))
  Planning Time: 0.811 ms
  Execution Time: 1898.255 ms
(4 rows)

```

Question 12

Observations

1. Creating an composite index (pid, skill) helped with query very effectively

- Postgres is performing Index only scan as the select clause is readily available in the index itself. Hence it is blazingly fast even for large number of records

size n of relation PersonSkill	No index	With index
10^4	0.776	0.025
10^5	6.859	0.027
10^6	31.71	0.028
10^7	360.053	0.028

Query Plans

For 10^4 records

No Index Query Plan

```

QUERY PLAN
-----
Seq Scan on personskill  (cost=0.00..184.21 rows=1 width=14) (actual time=0.016..0.786 rows=1 loops=1)
  Filter: ((pid = 101) AND (skill = 'AI'::text))
  Rows Removed by Filter: 9013
Planning Time: 0.094 ms
Execution Time: 0.797 ms
(5 rows)

```

With Index Query Plan

```

QUERY PLAN
-----
Index Only Scan using personskill_pid_skill_idx on personskill  (cost=0.29..4.30 rows=1 width=14) (actual time=0.046..0.047 rows=1 loops=1)
  Index Cond: ((pid = 101) AND (skill = 'AI'::text))
  Heap Fetches: 0
Planning Time: 0.234 ms
Execution Time: 0.058 ms
(5 rows)

```

For 10^5 records

No Index Query Plan

```

QUERY PLAN
-----
Seq Scan on personskill  (cost=0.00..1847.31 rows=1 width=14) (actual time=0.017..7.034 rows=1 loops=1)
  Filter: ((pid = 102) AND (skill = 'AI'::text))
  Rows Removed by Filter: 90553
Planning Time: 0.053 ms
Execution Time: 7.047 ms
(5 rows)

```

With Index Query Plan

```

QUERY PLAN
-----
Index Only Scan using personskill_pid_skill_idx on personskill  (cost=0.42..4.44 rows=1 width=14) (actual time=0.063..0.064 rows=1 loops=1)
  Index Cond: ((pid = 102) AND (skill = 'AI'::text))
  Heap Fetches: 0
Planning Time: 0.335 ms
Execution Time: 0.076 ms
(5 rows)

```

For 10⁶ records

No Index Query Plan

```
QUERY PLAN
-----
Gather  (cost=1000.00..11557.49 rows=1 width=14) (actual time=0.183..31.747 rows=1 loops=1)
  Workers Planned: 2
  Workers Launched: 2
  -> Parallel Seq Scan on personskill  (cost=0.00..10557.39 rows=1 width=14) (actual time=15.999..25.964 rows=0 loops=3)
        Filter: ((pid = 160) AND (skill = 'AI'::text))
        Rows Removed by Filter: 302154
Planning Time: 0.048 ms
Execution Time: 31.761 ms
(8 rows)
```

With Index Query Plan

```
QUERY PLAN
-----
Index Only Scan using personskill_pid_skill_idx on personskill  (cost=0.42..4.44 rows=1 width=14) (actual time=0.055..0.056 rows=1 loops=1)
  Index Cond: ((pid = 160) AND (skill = 'AI'::text))
  Heap Fetches: 0
Planning Time: 0.254 ms
Execution Time: 0.068 ms
(5 rows)
```

For 10⁷ records

No Index Query Plan

```
QUERY PLAN
-----
Gather  (cost=1000.00..106540.20 rows=1 width=14) (actual time=0.708..366.035 rows=1 loops=1)
  Workers Planned: 2
  Workers Launched: 2
  -> Parallel Seq Scan on personskill  (cost=0.00..105540.10 rows=1 width=14) (actual time=233.990..354.330 rows=0 loops=3)
        Filter: ((pid = 120) AND (skill = 'AI'::text))
        Rows Removed by Filter: 3020812
Planning Time: 0.049 ms
Execution Time: 366.052 ms
(8 rows)
```

With Index Query Plan

```
QUERY PLAN
-----
Index Only Scan using personskill_pid_skill_idx on personskill  (cost=0.43..4.45 rows=1 width=14) (actual time=0.012..0.013 rows=1 loops=1)
  Index Cond: ((pid = 100) AND (skill = 'AI'::text))
  Heap Fetches: 0
Planning Time: 0.055 ms
Execution Time: 0.025 ms
(5 rows)
```

Question 13

Semi Join

Observations

- 1.

size n of relation PersonSkill	No index	With index
10^4	4.021	3.262
10^5	38.645	32.764
10^6	523.410	457.729
10^7	9592.538	8720.626

Query Plans

For 10^4 records

No Index Query Plan

```

QUERY PLAN
-----
Hash Semi Join  (cost=185.39..399.72 rows=2896 width=8) (actual time=1.415..3.828 rows=1843 loops=1)
  Hash Cond: (person.pid = personskill.pid)
    -> Seq Scan on person  (cost=0.00..145.00 rows=10000 width=8) (actual time=0.012..0.967 rows=10000 loops=1)
    -> Hash  (cost=162.35..162.35 rows=1843 width=4) (actual time=1.395..1.396 rows=1843 loops=1)
          Buckets: 2048  Batches: 1  Memory Usage: 81kB
          -> Seq Scan on personskill  (cost=0.00..162.35 rows=1843 width=4) (actual time=0.008..1.074 rows=1843 loops=1)
                Filter: (skill = 'AI'::text)
                Rows Removed by Filter: 7225
Planning Time: 0.373 ms
Execution Time: 3.973 ms
(10 rows)

```

With Index Query Plan

```

QUERY PLAN
-----
Hash Semi Join  (cost=87.57..301.90 rows=2896 width=8) (actual time=0.584..3.335 rows=1843 loops=1)
  Hash Cond: (person.pid = personskill.pid)
    -> Seq Scan on person  (cost=0.00..145.00 rows=10000 width=8) (actual time=0.009..1.032 rows=10000 loops=1)
    -> Hash  (cost=64.54..64.54 rows=1843 width=4) (actual time=0.566..0.567 rows=1843 loops=1)
          Buckets: 2048  Batches: 1  Memory Usage: 81kB
          -> Index Only Scan using personskill_skill_pid_idx on personskill  (cost=0.29..64.54 rows=1843 width=4) (actual time=0.023..0.303 rows=1843 loops=1)
                Index Cond: (skill = 'AI'::text)
                Heap Fetches: 0
Planning Time: 0.341 ms
Execution Time: 3.548 ms
(10 rows)

```

For 10^5 records

No Index Query Plan

```

QUERY PLAN
-----
Hash Semi Join  (cost=1849.03..3953.87 rows=32270 width=8) (actual time=12.274..37.028 rows=18119 loops=1)
  Hash Cond: (person.pid = personskill.pid)
    -> Seq Scan on person  (cost=0.00..1443.00 rows=100000 width=8) (actual time=0.009..8.995 rows=100000 loops=1)
    -> Hash  (cost=1623.28..1623.28 rows=18060 width=4) (actual time=12.244..12.245 rows=18119 loops=1)
          Buckets: 32768  Batches: 1  Memory Usage: 893kB
          -> Seq Scan on personskill  (cost=0.00..1623.28 rows=18060 width=4) (actual time=0.018..9.072 rows=18119 loops=1)
                Filter: (skill = 'AI'::text)
                Rows Removed by Filter: 72543
Planning Time: 0.306 ms
Execution Time: 38.719 ms
(10 rows)

```

With Index Query Plan

```

QUERY PLAN
-----
Hash Semi Join  (cost=826.22..2931.06 rows=32270 width=8) (actual time=6.257..31.265 rows=18119 loops=1)

```

```

Hash Cond: (person.pid = personskill.pid)
-> Seq Scan on person (cost=0.00..1443.00 rows=100000 width=8) (actual time=0.006..8.948 rows=100000 loops=1)
-> Hash (cost=600.47..600.47 rows=18060 width=4) (actual time=6.233..6.234 rows=18119 loops=1)
    Buckets: 32768 Batches: 1 Memory Usage: 893kB
    -> Index Only Scan using personskill_skill_pid_idx on personskill (cost=0.42..600.47 rows=18060 width=4) (actual time=0.071..3.205 rows=18119 loops=1)
        Index Cond: (skill = 'AI'::text)
        Heap Fetches: 0
Planning Time: 0.704 ms
Execution Time: 33.093 ms
(10 rows)

```

For 10^6 records

No Index Query Plan

```

QUERY PLAN
-----
Hash Semi Join (cost=19085.86..48736.60 rows=331534 width=8) (actual time=117.665..538.155 rows=180988 loops=1)
  Hash Cond: (person.pid = personskill.pid)
  -> Seq Scan on person (cost=0.00..14425.00 rows=1000000 width=8) (actual time=0.008..83.184 rows=1000000 loops=1)
  -> Hash (cost=16215.64..16215.64 rows=174898 width=4) (actual time=117.446..117.447 rows=181094 loops=1)
      Buckets: 131072 Batches: 4 Memory Usage: 2625kB
      -> Seq Scan on personskill (cost=0.00..16215.64 rows=174898 width=4) (actual time=0.007..82.050 rows=181094 loops=1)
          Filter: (skill = 'AI'::text)
          Rows Removed by Filter: 724957
Planning Time: 0.277 ms
Execution Time: 548.909 ms
(10 rows)

```

With Index Query Plan

```

QUERY PLAN
-----
Hash Semi Join (cost=8709.55..38478.98 rows=340641 width=8) (actual time=62.806..444.022 rows=181151 loops=1)
  Hash Cond: (person.pid = personskill.pid)
  -> Seq Scan on person (cost=0.00..14425.00 rows=1000000 width=8) (actual time=0.012..81.061 rows=1000000 loops=1)
  -> Hash (cost=5813.24..5813.24 rows=176504 width=4) (actual time=62.505..62.507 rows=181054 loops=1)
      Buckets: 131072 Batches: 4 Memory Usage: 2627kB
      -> Index Only Scan using personskill_skill_pid_idx on personskill (cost=0.42..5813.24 rows=176504 width=4) (actual time=0.028..27.646 rows=181054 loops=1)
          Index Cond: (skill = 'AI'::text)
          Heap Fetches: 0
Planning Time: 0.208 ms
Execution Time: 454.270 ms
(10 rows)

```

For 10^7 records

No Index Query Plan

```

QUERY PLAN
-----
Hash Semi Join (cost=191842.00..493738.08 rows=3734414 width=8) (actual time=2004.333..9675.957 rows=1811473 loops=1)
  Hash Cond: (person.pid = personskill.pid)
  -> Seq Scan on person (cost=0.00..144247.77 rows=9999977 width=8) (actual time=0.020..1577.721 rows=10000000 loops=1)
  -> Hash (cost=162196.96..162196.96 rows=1806883 width=4) (actual time=2001.079..2001.080 rows=1812756 loops=1)
      Buckets: 131072 Batches: 32 Memory Usage: 3013kB
      -> Seq Scan on personskill (cost=0.00..162196.96 rows=1806883 width=4) (actual time=0.013..1440.731 rows=1812756 loops=1)
          Filter: (skill = 'AI'::text)
          Rows Removed by Filter: 7250716
Planning Time: 0.455 ms
Execution Time: 9840.218 ms
(10 rows)

```

With Index Query Plan

```

QUERY PLAN
-----
Hash Semi Join (cost=89100.05..390998.23 rows=3734560 width=8) (actual time=1037.949..8629.806 rows=1811473 loops=1)

```

```
Hash Cond: (person.pid = personskill.pid)
-> Seq Scan on person (cost=0.00..144248.00 rows=10000000 width=8) (actual time=0.022..1516.697 rows=10000000 loops=1)
-> Hash (cost=59454.13..59454.13 rows=1806954 width=4) (actual time=1034.911..1034.912 rows=1812756 loops=1)
    Buckets: 131072 Batches: 32 Memory Usage: 3013kB
    -> Index Only Scan using personskill_skill_pid_idx on personskill (cost=0.43..59454.13 rows=1806954 width=4) (actual time=0.780..505.690 rows=1812756 loops=1)
        Index Cond: (skill = 'AI'::text)
        Heap Fetches: 0
Planning Time: 2.328 ms
Execution Time: 8798.384 ms
(10 rows)
```

Anti Semi Join

Observations

1.

size n of relation PersonSkill	No index	With index
10^4		3.262
10^5	38.645	32.764
10^6	523.410	457.729
10^7	9592.538	8720.626

Query Plans

For 10^4 records

No Index Query Plan

```
QUERY PLAN
-----
Seq Scan on person (cost=166.97..336.97 rows=5000 width=8) (actual time=1.256..3.182 rows=8112 loops=1)
  Filter: (NOT (hashed SubPlan 1))
  Rows Removed by Filter: 1888
  SubPlan 1
    -> Seq Scan on personskill (cost=0.00..162.35 rows=1847 width=4) (actual time=0.007..0.905 rows=1847 loops=1)
        Filter: (skill = 'AI'::text)
        Rows Removed by Filter: 7221
Planning Time: 0.114 ms
Execution Time: 3.731 ms
```

With Index Query Plan

For 10^5 records

No Index Query Plan

With Index Query Plan

For 10^6 records

No Index Query Plan

With Index Query Plan

For 10^7 records

No Index Query Plan

With Index Query Plan